

I am a micro labor economist studying how organizational decisions and social institutions shape hiring and labor-market opportunities. My job market paper shows that employers' decisions about when to begin contacting applicants change which candidates receive consideration, even when qualifications are held fixed. This finding motivates my broader research on how workers reach employers and how the organization of work shapes employers' hiring needs and constraints. I study how training systems and employer networks connect workers to jobs, how professionals learn from feedback, and how organizations allocate work across competing demands. I build evidence on otherwise unobserved decision processes through experiments and linked administrative, historical, and text-based records. I combine computational measurement, causal research designs, and economic models to distinguish mechanisms and quantify their consequences.

Access to work and the formation of matches

My job market paper, *Time Is of the Essence*, studies the role of job application timing in recruitment. It identifies a firm-side source of the return to applying early: employers begin contacting candidates before their applicant pools are complete. I repurpose the correspondence experiment of Kline, Rose, and Walters (2022), which sent more than 83,000 applications to over 11,000 entry-level vacancies at 108 large U.S. firms. Its résumé randomization and controlled submission sequence allow me to study how firms' contact decisions change as applications arrive: within each vacancy, applications in the first half of the experimental sequence are on average 1.16 percentage points more likely to receive employer contact within 30 days. This effect is equivalent to 57 percent of the Black–White contact gap in the same setting.

The mechanism is a threshold decision. A higher predicted screening value for the best application received so far prompts earlier contact initiation, defined as the first observed experimental contact for a vacancy. Using variation from randomized résumé attributes to instrument whether contact has begun, I find that contact initiation reduces consideration of subsequent arrivals. To quantify firms' tradeoff between moving quickly and waiting for a larger applicant pool, I develop and estimate a sequential recruiting model that separates applicant arrival intensity from the cost of delay. Conditional on experimental contact initiation, the model implies that in 52 percent of cases, the highest-ranked experimental application has not yet arrived when the vacancy enters the contact phase. The average shortfall due to this timing friction equals 6.6 percent of the initiating application's expected screening value. Holding the contact initiation date fixed, better screening of the current pool cannot eliminate this shortfall. The results have important welfare implications: firms compete for workers partly through how quickly they initiate contact. For applicants, the timing premium can induce costly competition in speed without improving match quality, while disadvantaging those less able to apply quickly. Steeper timing gradients in tighter labor markets connect this concrete recruiting margin to the cyclicity of matching efficiency. Together, my job market paper identifies and quantifies a within-vacancy timing friction, providing an empirical foundation for search models that explain who receives consideration and how matches are made.

Workers' ability to respond to changing hiring needs depends on access to training. In *From AI Exposure to Training Access*, I study how funding and delivery capacity shape public training

systems' response to changing demands in the AI era. I link 2017–2024 Workforce Innovation and Opportunity Act (WIOA) records to program listings and state funding. For each local workforce board, I measure participants' occupational AI exposure and the board's prior experience delivering AI-related training. After ChatGPT's launch, the training participation rate rose more at boards with greater AI exposure and prior AI-training experience. This increase was larger where states received larger federal workforce allocations. Yet the increase in participation was not accompanied by a clear shift toward AI-related fields, even as more such programs became eligible for public funding. Training also became more concentrated in these boards relative to their shares of statewide clients. These findings connect training access to public finance and institutional capacity, raising a distributional question about who can benefit from established training systems.

Networks also shape which employers job seekers can reach. In *Employer Networks as Organizational Capital in Professional Education*, I link records of U.S. law-school reorganizations, closures, and openings to graduates' career trajectories to trace changes in employer relationships. After reorganizations, hiring shifts toward employers connected to the new parent or stronger merger partner, while other inherited ties weaken. After closures, employers do not replace the lost pipeline by hiring more graduates from other law schools. New schools gain access to employers linked to parents or sponsors without commensurate growth in those employers' total law-graduate hiring. These patterns suggest that professional schools maintain employer networks as organizational capital: difficult to expand, selectively recombined, and incompletely absorbed after closure. This extends labor-market network analysis from individual contacts to institutional relationships.

Cultural norms influence workers' aspirations and the choices open to them. In *Beyond the Altar*, with Monia Tomasella and Juan David Torres, I ask whether Catholic convents changed women's family and work choices in the United States, 1850–1940. Convents offered a respected alternative to marriage, combining material support with roles in education and care. We are linking religious-order arrivals reconstructed from historical catalogs to women's census records from childhood through adulthood to study how exposure at different life stages affects marriage and labor-market participation. The project asks how visible role models, material opportunities outside marriage, and expectations about gender roles interact to shape women's career choices.

On the employer side, how organizations deploy their existing workforce shapes their recruiting needs. In *Reallocating Expertise in the AI Era*, I link patent applications, examiner assignments, and examination records at the U.S. Patent Office to study how specialized capacity expands as demand shifts. A backfill pattern emerges: as AI-related applications increase, experienced specialists shift toward them and away from other work, while newly observed examiners account for large increases in non-AI examination. AI cases show improved reference coverage in the initial examination and less later rework relative to non-AI cases. These patterns connect the expansion of specialized capacity to staffing elsewhere and raise the question: when does hiring complement the reallocation of the existing workforce? This motivates a closer look at how review and delegation shape professional decisions and the allocation of work within organizations.

Learning and coordination within organizations

In *When Hierarchies Teach*, with Daniel L. Chen, I study two ingredients of hierarchical review's teaching function: the information in a correction and its recipient's responsiveness. In U.S. asylum courts, feedback follows a denial, an appeal, and review by the Board of Immigration Appeals (BIA). Using court and appellate records from 1986–2019, we distinguish surprise from expected

reversals with a benchmark built from information available before appellate decisions. Both overturn denials, but surprise reversals are followed by a 1.57-percentage-point larger first-week increase in grant rates. Among judges whose grant rates respond more strongly, reversal rates among denied cases vary less with leniency, with the clearest contrast among stricter judges. Our threshold model interprets this pattern as closer distributional alignment with the BIA standard and connects responsiveness to attention and effective updating. Together, this paper explains how a common hierarchy can generate differences in learning and alignment across its members. In *Strategic Predictability*, I extend this learning agenda with a dynamic model that separates learning about one's own payoffs from shaping an opponent's expectations. I test its implications using professional tennis serves, examining how creating predictable patterns can shape what others learn to expect.

In *Delegation under Deadlines*, I use patent prosecution to study how task competition constrains principals' ability to obtain timely legal services. Patent lawyers serve multiple clients, each with several applications and response deadlines. I am linking application histories to lawyers and clients to reconstruct completed responses and pending work. Holding the number of approaching deadlines fixed, I ask whether an application's legal team is less likely to respond when more competing deadlines belong to the same client rather than others. Comparisons across client relationships and between companies' own legal teams and outside law firms connect these priorities to organizational form. This project brings a principal-agent perspective to the study of task management and the organization of work.

Future directions

Two planned projects bring these organizational questions back to hiring and mobility. The first asks how the organization of recruiting shapes learning about applicants and who receives consideration. My JMP documents two motivating patterns. Earlier contacts predict later contact decisions differently across applicant groups. Timing gradients are also steeper at firms using fewer distinct callback phone numbers per vacancy receiving a callback, a proxy for more centralized recruiting. I plan a large survey of HR managers to collect field data on recruiting task assignments, decision authority, and team management. This will guide experiments independently varying information about earlier applicants and recruiters' organizational arrangements. Holding later applications fixed, I will test how these arrangements change responses to earlier information within and across applicant groups. A sequential model will quantify how recruiting organization affects belief updating and the tradeoff between recruiting speed and breadth of applicant consideration.

A second project asks when promotion bottlenecks redirect experienced workers toward other employers. I will reconstruct firm-specific job ladders from career histories and measure next-level openings relative to candidates identified by prior roles and trajectories. Linking these measures to job postings, I will compare workers at similar career stages and ask whether those with fewer internal promotion opportunities are more likely to move into more senior jobs elsewhere. The project connects firms' promotion structures to the supply of experienced workers available to other employers, extending the study of job ladders across organizational boundaries.

My research connects the organization of work to the distribution of labor-market opportunities. By making the decisions behind hiring, learning, and work allocation measurable and modeling their consequences, I study how institutions shape not only who works where, but what workers and organizations can accomplish together.