

I am a micro labor economist studying how organizational decisions and social institutions shape hiring and labor-market opportunities. My job market paper shows that employers' decisions about when to begin contacting applicants change which candidates receive consideration, even when qualifications are held fixed. This finding motivates my broader research on how workers reach employers and how the organization of work shapes employers' hiring needs and constraints. I study how training systems and employer networks connect workers to jobs, how professionals learn from feedback, and how organizations allocate work across competing demands. I build evidence on otherwise unobserved decision processes through experiments and linked administrative, historical, and text-based records. I combine computational measurement, causal research designs, and economic models to distinguish mechanisms and quantify their consequences.

Access to work and the formation of matches

My job market paper, *Time Is of the Essence*, studies the role of job application timing in recruitment. It identifies a firm-side source of the return to applying early: employers begin contacting candidates before their applicant pools are complete. I repurpose the correspondence experiment of Kline, Rose, and Walters (2022), which sent more than 83,000 applications to over 11,000 entry-level vacancies at 108 large U.S. firms. Its résumé randomization and controlled submission sequence allow me to study how firms' contact decisions change as applications arrive: within each vacancy, applications in the first half of the experimental sequence are on average 1.16 percentage points more likely to receive employer contact within 30 days. This effect is equivalent to 57 percent of the Black–White contact gap in the same setting.

The mechanism is a threshold decision. A higher predicted screening value for the best application received so far prompts earlier contact initiation, defined as the first observed experimental contact for a vacancy. Using variation from randomized résumé attributes to instrument whether contact has begun, I find that contact initiation reduces consideration of subsequent arrivals. To quantify firms' tradeoff between moving quickly and waiting for a larger applicant pool, I develop and estimate a sequential recruiting model that separates applicant arrival intensity from the cost of delay. Conditional on experimental contact initiation, the model implies that in 52 percent of cases, the highest-ranked experimental application has not yet arrived when the vacancy enters the contact phase. The average shortfall due to this timing friction equals 6.6 percent of the initiating application's expected screening value. Holding the contact initiation date fixed, better screening of the current pool cannot eliminate this shortfall. The results have important welfare implications: firms compete for workers partly through how quickly they initiate contact. For applicants, the timing premium can induce costly competition in speed without improving match quality, while disadvantaging those less able to apply quickly. Steeper timing gradients in tighter labor markets connect this concrete recruiting margin to the cyclicity of matching efficiency. Together, my job market paper identifies and quantifies a within-vacancy timing friction, providing an empirical foundation for search models that explain who receives consideration and how matches are made.

Workers' ability to respond to changing hiring needs depends on access to training. In *From AI Exposure to Training Access*, I study how funding and delivery experience shape public training

systems' response to AI. I link 2017–2024 Workforce Innovation and Opportunity Act records to program listings and state funding, measuring AI exposure from clients' preexisting occupations and comparing boards within the same state and year. After ChatGPT's launch, the share of clients receiving training rose more at boards with greater AI exposure and prior AI-training experience. This increase was larger where states received larger federal workforce allocations. Yet the increase in training participation was not accompanied by a clear shift toward AI-related fields, even as more such programs became eligible for public funding. Training also became more concentrated in these boards relative to their shares of statewide clients. These findings connect retraining access to public finance and institutional capacity, raising a distributional question about who can benefit from established training systems.

Job seekers' access to employers also depends on networks maintained by institutions. In *Employer Networks as Organizational Capital in Professional Education*, I link U.S. law-school histories to graduates' careers to study how these networks survive organizational change. After reorganizations, hiring shifts toward employers connected to the new parent or stronger merger partner, while other inherited ties weaken. After closures, hiring from other schools does not replace the lost pipeline. New schools gain access to employers linked to parents or sponsors without commensurate growth in those employers' total hiring of law graduates. These patterns suggest organizational capital that is difficult to expand, selectively recombined, and incompletely absorbed after closure, extending network analysis from individual contacts to relationships maintained by institutions.

Social norms also shape job seekers' preferences and constraints. In *Beyond the Altar*, with Monia Tomasella and Juan David Torres, I ask whether Catholic convents changed family and work choices among women who never entered religious life. In the United States, convents combined material support outside marriage with respected roles for women in education and care. We have reconstructed religious-order arrivals from historical catalogs and assembled linked census records. We are developing an analysis following girls from childhood locations to adult marriage and occupations, distinguishing individual choices from changes in county populations. The project asks how material opportunities and expectations about women's roles interact to shape their choices.

On the employer side, how organizations deploy their existing workforce shapes whom they need to recruit. In *Reallocating Expertise in the AI Era*, I link patent applications, examiner assignments, and written decisions at the U.S. Patent Office to study how specialized capacity expands across workers and tasks. As AI-related applications increase, experienced specialists shift toward them and away from other work, while newly observed examiners contribute especially large increases in non-AI examination. I also examine decision content and subsequent rework. These patterns raise a broader question: how does the organization of expertise shape recruiting needs? My second strand examines the review and delegation arrangements governing how that expertise is used.

Learning and coordination within organizations

In *When Hierarchies Teach*, with Daniel L. Chen, I study two ingredients of hierarchical review's teaching function: the information in a correction and its recipient's responsiveness. Using U.S. asylum-court and appellate records from 1986–2019, we study feedback generated by denials, appeals, and Board of Immigration Appeals decisions. A benchmark built from information available before each appellate decision distinguishes expected from surprise reversals. Both overturn denials, but surprise reversals are followed by a 1.57-percentage-point larger first-week

increase in grant rates. Among judges whose grant rates respond more strongly, reversal rates among denied cases vary less with leniency, with the clearest contrast among stricter judges. Our threshold model interprets this pattern as closer distributional alignment with the Board's standard and connects responsiveness to attention and effective updating. The paper separates the information supplied by review from how recipients use it to align professional decisions.

In *Delegation under Deadlines*, I ask how delegation incentives shape work organization and how task competition constrains principals' ability to obtain timely service. Patent lawyers serve multiple clients, each with several applications and response deadlines. I am linking application histories to lawyers and clients to reconstruct completed responses and pending work. Holding the number of approaching deadlines fixed, I ask whether an application's legal team is less likely to respond when more competing deadlines belong to the same client rather than others. Comparisons across client relationships and between in-house and outside counsel connect these priorities to organizational form.

I also develop a dynamic model of strategic predictability and test its choice implications using professional tennis serves. The model separates learning about one's own payoffs from shaping an opponent's expectations. Predictable patterns can raise future payoffs even when rarely broken; fitted serve policies are consistent with this motive. This extends my learning agenda to how people shape the expectations against which they act.

Future directions

Two planned projects bring these organizational questions back to hiring and mobility. The first asks how the organization of recruiting shapes learning about applicants and who receives consideration. The JMP documents group-specific dependence of contact decisions on earlier contacts. Timing gradients are also steeper at firms using fewer distinct callback phone numbers per vacancy receiving a callback, a proxy for centralization. I plan a large survey of HR managers to map task assignments, decision authority, and collaboration. This will guide experiments independently varying information about earlier applicants and recruiters' organizational arrangements. Holding later applications fixed, I will test how these arrangements change responses to earlier information within and across applicant groups. A sequential model will quantify how these arrangements change belief updating and the tradeoff between recruiting speed and breadth of consideration. A later extension will examine how job seekers change their application decisions when informed about recruiting structures.

A second project asks when promotion bottlenecks redirect experienced workers toward other employers. I will reconstruct firm-specific job ladders from career histories and measure next-level openings relative to candidates identified by prior roles and trajectories. Linking these measures to job postings, I will compare moves into more senior roles elsewhere among similarly positioned workers facing different internal advancement constraints and outside opportunities. The project connects firms' promotion structures to the supply of experienced workers available to other employers, extending the study of job ladders across organizational boundaries.

My research connects the organization of work to the distribution of labor-market opportunities. By making the decisions behind hiring, learning, and work allocation measurable and modeling their consequences, I study how institutions shape not only who works where, but what workers and organizations can accomplish together.